

Luowei Zhou

1301 Beal Ave., EECS Building 4338, Ann Arbor, MI 48105, USA
luoweizhou.github.io | luozhou@umich.edu | (734)757-3282

EDUCATION

University of Michigan

Ann Arbor, Michigan, USA

Ph.D. program in Robotics (Computer Vision)

Sept. 2015 – Dec. 2019 (expected)

- **Research Interests:** Vision and language, video understanding, deep learning
- **Courses:** Advanced Computer Vision, Natural Language Processing, Machine Learning, Optimization
- **Academics:** Curriculum GPA: **4.00/4.00**

Nanjing University

Nanjing, Jiangsu, China

Bachelor of Engineering in Automation

Sept. 2011 – Jun. 2015

- **Courses:** Computer Vision, Artificial Intelligence, Advanced Programming Language, Data Structure
- **Academics:** Overall GPA: **91.8/100**, Major GPA: **93.0/100**

PREPRINTS AND PUBLICATIONS

L. Zhou, Y. Kalantidis, X. Chen, J. J. Corso, and M. Rohrbach, “*Grounded Video Description*”, in submission.

L. Zhou, Y. Zhou, J. J. Corso, R. Socher, and C. Xiong, “*End-to-End Dense Video Captioning with Masked Transformer*”, CVPR 2018. (**spotlight**) AR: 9%; h5: 158

L. Zhou, N. Louis, and J. J. Corso, “*Weakly-Supervised Video Object Grounding from Text by Loss Weighting and Object Interaction*”, BMVC 2018. AR: 30%; h5: 42

L. Zhou, C. Xu, and J. J. Corso, “*Towards Automatic Learning of Procedures from Web Instructional Videos*”, AAAI 2018. (**oral**) AR: 11%; h5: 56

L. Zhou, C. Xu, P. Koch, and J. J. Corso, “*Watch What You Just Said: Image Captioning with Text-Conditional Attention*”, ACM Multimedia (Thematic Workshops) 2017: 305-313. (**pitch**)

L. Zhou, P. Yang, C. Chen, and Y. Gao, “*Multi-agent Reinforcement Learning with Sparse Interactions by Negotiation and Knowledge Transfer*”, IEEE Transactions on Cybernetics 2017, 47 (5): 1238 - 1250. SCI IF: 7.38; h5: 73

L. Zhou, P. Yang, and C. Chen, “*Multi-agent Reinforcement Learning with Sparse Interactions by Negotiation and Knowledge Transfer*”, IJCAI (Workshops) 2016. (**oral**)

L. Zhou, Y. Shi, J. Wang, and P. Yang, “*A Balanced Heuristic Mechanism for Multi-robot Task Allocation of Intelligent Warehouses*”, Mathematical Problems in Engineering 2014: 1–10. SCI IF: 0.80; h5: 39

WORK EXPERIENCE

Facebook AI Research (FAIR)

Menlo Park, CA, USA

Research Intern with Marcus Rohrbach, Yannis Kalantidis, and Xinlei Chen

May 2018 – Present

University of Michigan, EECS

Ann Arbor, MI, USA

Graduate Research Assistant with Prof. Jason Corso

Apr. 2016 – Present

Salesforce Research (Metamind)

Palo Alto, CA, USA

Deep Learning Research Intern with Caiming Xiong and Richard Socher

May 2017 – Aug. 2017

PROFESSIONAL ACTIVITIES

Co-organizer, CVPR 2018 Workshop on Fine-grained Instructional Video Understanding, with Jason Corso, Josef Sivic and Ivan Laptev

Reviewer, CVPR 2019, TPAMI 2018&2017, IJCV 2018, ITS 2018&2017, ICRA 2017, NIPS 2016, CVIU 2016

Volunteer, RSS 2016

Attendee, CVPR 2018&2016, AAAI 2018, BMVC 2018, ACM Multimedia 2017, IJCAI 2016

HONORS AND AWARDS

Outstanding Winner Awards (0.2%), Mathematical Contest in Modeling (MCM) 2013

Sienhua New and Tsu Way Shen Memorial Award (Top 1), of University of Michigan 2015

Best Undergrad Thesis (Top 1), of Jiangsu Province 2015

National Scholarship (1%), of Nanjing University 2012

Red Sun Scholarship, of Nanjing University 2014

Travel Grant, of University of Michigan 2016-2018

RESEARCH EXPERIENCE

Grounded Video Description

Facebook AI Research

Supervisors: Dr. Marcus Rohrbach, Dr. Yannis Kalantidis, and Dr. Xinlei Chen May 2018 – present

- Introduced a large-scale video description and grounding dataset, called ActivityNet-Entities, where we annotated noun phrases (& objects) from sentence descriptions in videos as spatial bounding boxes. ActivityNet-Entities contains over 158k labelled boxes for 52k video clips.
- Proposed a unified framework for video and image description, where a grounding module dynamically detects objects in the scene and provides visual clues to the captioning module.
- Set new state-of-the-art performance on video description and image description and demonstrated that our generated sentences are more explainable through grounding.

Fine-grained Instructional Video Understanding

University of Michigan

Supervisor: Prof. Jason Corso

Sept. 2016 – present

- Introduced [YouCook2](#) dataset, which contains temporally localized recipe sentence annotations and bounding boxes for 2000 YouTube cooking videos.
- Tackled a series of problems related to instructional video understanding: i) event proposal (AAAI 2018), ii) dense video captioning (CVPR 2018), iii) weakly supervised object grounding from language description (BMVC 2018).
- *Event proposal*: Proposed an event proposal and sequential modeling network that can temporally localize procedure steps in web instructional videos and capture the temporal structure of the video.
- *Dense video captioning*: Caption generation for event proposals. See Page 3 for more details.
- *Weakly supervised object grounding*: Given a video and the corresponding description, localize the objects mentioned from the description in the video as bounding boxes. No box is given for training.
- Current focuses: Action-conditioned object grounding, modeling of object state transitions, multi-task learning of object grounding, action recognition, and captioning on videos.

Dense-Captioning Events in Video and Temporal Action Proposal

Supervisors: Dr. Caiming Xiong and Dr. Richard Socher

Salesforce Research
May 2017 – Aug. 2017

- Introduced a Self-attention-based video captioning model and improved our previously proposed action/event proposal network with carefully-designed Temporal Convolutional Networks.
- Proposed to bridge event proposal and captioning by a differentiable visual mask and achieved state-of-the-art results on dense video captioning.

Text-conditional Visual Captioning with Guiding LSTM

Supervisor: Prof. Jason Corso

University of Michigan
Mar. 2016 – Nov. 2016

- Proposed an encoder-decoder image captioning method through explicit text-conditional image guidance.
- Extended the work to video captioning by leveraging audio features for the extra guidance.

End-to-End Grasping with Deep Reinforcement Learning

Supervisor: Prof. Satinder Singh

University of Michigan
Sept. 2015 – Apr. 2016

- Applied state-of-the-art Deep RL algorithm named Deep Q-network (DQN) to robot grasping tasks.
- Built an API between physics engine MuJoCo and the DQN module.

Research on Multi-Agent Reinforcement Learning with Sparse Interactions

Supervisors: Prof. Chunlin Chen, Dr. Pei Yang, Dr. Yang Gao

Nanjing University
Dec. 2014 – Jul. 2015

- Introduced the concept of equilibrium into traditional sparse-interaction-based MARL algorithms and proposed a knowledge transfer approach to initialize the joint-state Q table.
- Applied the proposed algorithm in a real-world setting, i.e., our intelligent warehouse simulator.

Multi-Robot Task Allocation and Path Planning in Dynamic Environments

Supervisor: Dr. Pei Yang

Nanjing University
Nov. 2013 – Jul. 2014

- Proposed a Balanced Heuristic Mechanism to balance task allocation in multi-robot systems.
- Built an intelligent warehouse simulator from scratch using C/OpenGL for model evaluation.

PROFICIENCY AND SKILLS

Technical Skills: PyTorch/Torch, Python, C/C++, Linux, Git, LaTeX, Matlab, Caffe, HTML, CSS, JS etc.
Languages: English (proficient) and Mandarin (native)

REFERENCES

Prof. Jason Corso, Associate Professor, University of Michigan, jjcorso@umich.edu

Prof. Chenliang Xu, Assistant Professor, University of Rochester, chenliang.xu@rochester.edu

Dr. Marcus Rohrbach, Research Scientist, Facebook AI Research, mrf@fb.com

Dr. Yannis Kalantidis, Research Scientist, Facebook Research, yannisk@fb.com

Dr. Xinlei Chen, Research Scientist, Facebook AI Research, xinleic@fb.com

Dr. Caiming Xiong, Director of Research, Salesforce Research, cxiong@salesforce.com